


K-nearest Neighbors




**Show me your
friends and I'll show
you your future.**




*K-nearest
neighbors*

**Show me your
~~friends~~ and I'll show
you your future.**


Example



Example

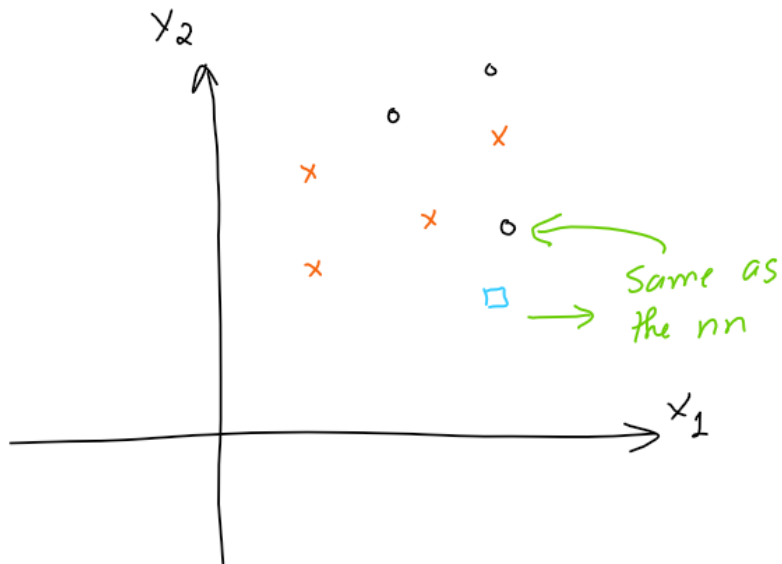


Example




Example

1 - nearest neighbor




Example



- 1NN model would classify the new sample as o

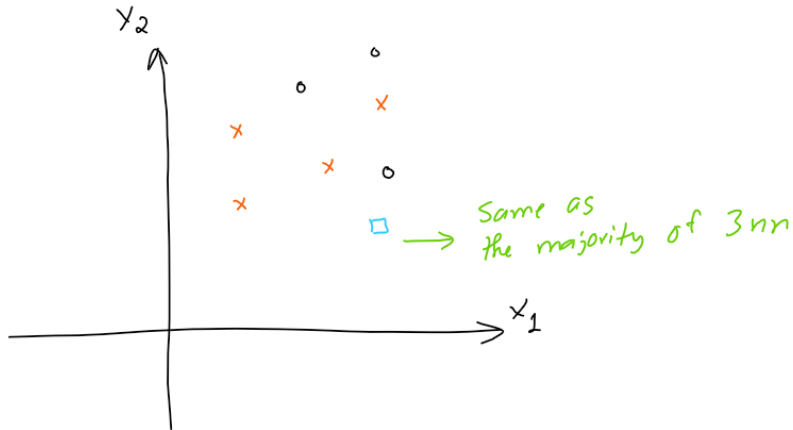
Example

3-nearest neighbor



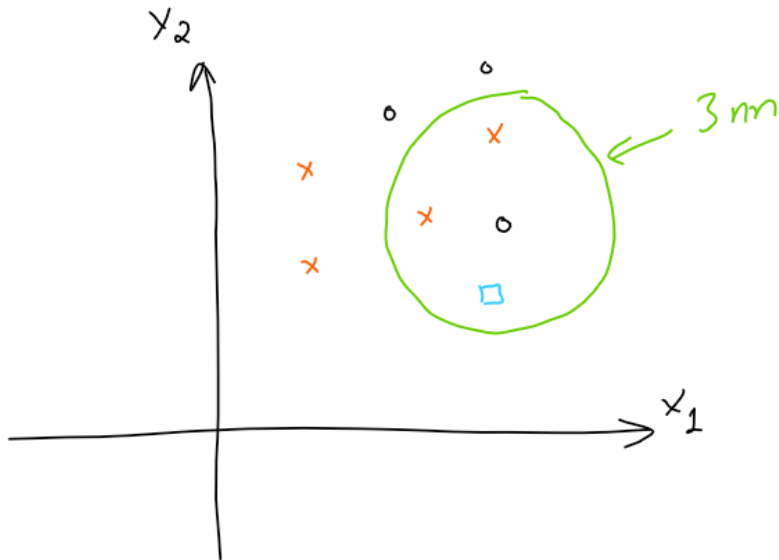
Example

3-nearest neighbor

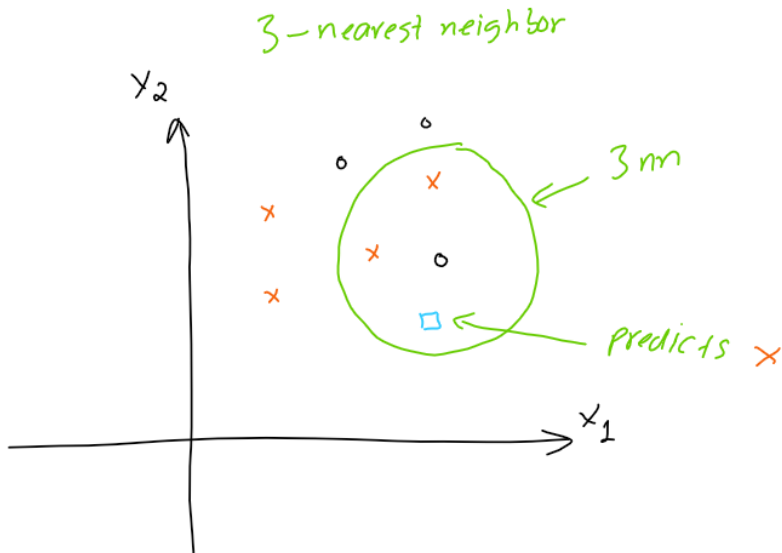


Example

3-nearest neighbor



Example



- 3NN model would classify the new sample as x .

KNN Algorithm

- ▶ **Prediction Rule:** Look at the K most similar training examples *categorical response*
- ▶ For **classification**: assign the *majority* class label (majority voting)
- ▶ For **regression**: assign the *average* response
continuous response

KNN Algorithm

- ▶ **Step 1:** Standardize the variables
- ▶ **Step 2:** Compute the test point's distance from each training point
- ▶ **Step 3:** Sort the distances in ascending (or descending) order
- ▶ **Step 4:** Use the sorted distances to select the K nearest neighbors
- ▶ **Step 5:** Use majority rule (for classification) or averaging (for regression)

Algorithm Requirements

- ▶ The algorithm requires:
 - ▶ **Tuning Parameter K**: number of nearest neighbors to look for
 - ▶ **Distance function**: To compute the similarities between examples

Distances

- ▶ The K-NN algorithm requires computing distances of the test example from each of the training examples
- ▶ Several ways to compute distances
- ▶ The choice depends on the type of the variables in the data

Distances

- ▶ **Euclidean distance** is commonly used when there are continuous variables

$$d(u, v) = \sqrt{(u_1 - v_1)^2 + (u_2 - v_2)^2 + \dots + (u_n - v_n)^2},$$

where

$$u = [u_1, u_2, \dots, u_n]$$

and

$$v = [v_1, v_2, \dots, v_n]$$

- ▶ **Hamming distance** is commonly used when there are categorical variables: $d(u, v) =$ Number of times u and v are different.

Distance Example

	Sex	Age	Class
u	Male	27	A
v	Female	30	A
Difference	1	3	0

$$d(u, v) = \sqrt{1^2 + 3^2 + 0^2} = \sqrt{10}$$

- Notice: Usually the **Age** variable needs to be standardized to have the range from 0 to 1 before calculating the distance.

Why standardizing the variables?

► Considering the following data:

data

person

	Age	Salary (\$1000)
A	15	90,000
B	30	80,000
C	80	87,000



We have

$$\underline{AB} = d(A, B) = \sqrt{(30 - 15)^2 + (80 - 90)^2} = 18.03$$

$$\underline{AC} = d(A, C) = \sqrt{(80 - 15)^2 + (87 - 90)^2} = 65.07$$

Thus,


$$\underline{AB} < \underline{AC}$$

B is closer to A than C

However, with the **same** data

Factor 2

	Age	Salary (\$)
A	15	90,000
B	30	89,000
C	80	87,000



We have

$$AB = d(A, B) = \sqrt{(30 - 15)^2 + (80,000 - 90,000)^2} = \underline{10,000.01}$$

$$AC = d(A, C) = \sqrt{(80 - 15)^2 + (87,000 - 90,000)^2} = \underline{3000.70}$$

Thus,

$$AB > AC$$

*C is closer to A
than B*


Not good!

Why standardizing the variables?

- ▶ Distances are affected by scalar-multiplication. Hence, the units of the variables will affect the distances.
- ▶ Standardizing variables will cancel this effect.

Why standardizing the variables?

- ▶ Common practice: Standardize variables to have the range from 0 to 1:


$$\text{Standardized } X = \frac{X - X_{\min}}{X_{\max} - X_{\min}}$$

Standardize the previous data (in either unit for salary):

	Age	Salary (\$1000)
A	0	1
B	0.23	0.7
C	1	0

rate 1
rate 2
⋮
←
←
B
C

We have

$$AB = d(A, B) = \sqrt{(0 - .23)^2 + (1 - .7)^2} = 0.38$$

$$AC = d(A, C) = \sqrt{(0 - 1)^2 + (1 - 0)^2} = 1.41$$

Choice of K - Neighborhood Size

- ▶ **Question:** Does larger or smaller K tend to overfit the model?

Choice of K - Neighborhood Size

- ▶ **Question:** Does larger or smaller K tend to overfit the model?
- ▶ **Hint:** Which one performs better on the training data?

Choice of K - Neighborhood Size

- ▶ Small K

- ▶ Creates many small regions for each class
- ▶ May lead to non-smooth decision boundaries and **overfit**


- ▶ Large K

- ▶ Creates fewer larger regions
- ▶ Usually leads to smoother decision boundaries (caution: too smooth decision boundary can **underfit**)

Weights in KNN

- If A, B, C are three nearest neighbors of D , then the predicted probability of the D by **3NN** is given by

$$\text{Predicted Probability of } D = \frac{w_A \cdot y_A + w_B \cdot y_B + w_C \cdot y_C}{w_A + w_B + w_C}$$



Weights in KNN

- ▶ **Uniform Weights:** All points in each neighborhood are weighted equally when predicting.
- ▶ If A , B , C are three nearest neighbors of D , then the predicted probability of the D by **3NN** with uniform **weights** becomes:

$$\text{Predicted Probability of } D = \frac{y_A + y_B + y_C}{3}$$

- ▶ Uniform weights are the default weights when using KNN

Weights in KNN

- ▶ **Distance Weights** weight points by the *inverse* of their distance. in this case, closer neighbors of a query point will have a greater influence than neighbors which are further away.
- ▶ If A , B , C are three nearest neighbors of D , then the predicted probability of the D by **3NN** with **distance weights** is:

$$\text{Predicted Probability of } D = \frac{\frac{1}{DA} \cdot y_A + \frac{1}{DB} \cdot y_B + \frac{1}{DC} \cdot y_C}{\frac{1}{DA} + \frac{1}{DB} + \frac{1}{DC}}$$

Weights in KNN



- **Uniform Weights:** All the three neighbors are weighted the same, so the majority vote predicts x .
- For all neighbors:

$$Weight = 1$$

Weights in KNN

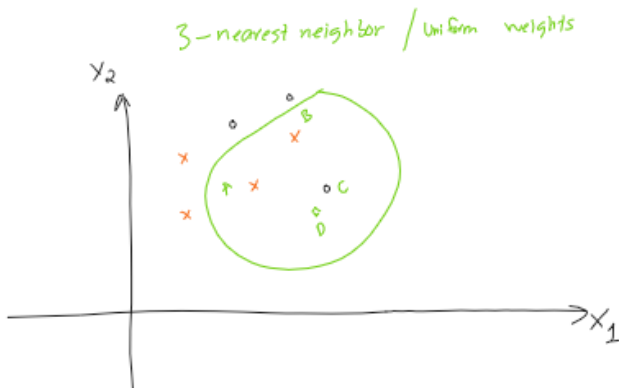


- **Distance Weights:** The closest neighbor (best friend!) is weighted more than the the two-further-away neighbors, so the weighted vote predicts o .

$$Weight = \frac{1}{Distance}$$


An Example of Uniform Weights

Use the uniform weights to calculate the predicted probability and the prediction of **3NN** for D. Consider x as 1 and o as 0.



Solution

3-nearest neighbor / Uniform weights



weight = 1


$$\text{predicted prob. of 0} = \frac{1 \cdot y_A + 1 \cdot y_B + 1 \cdot y_C}{3}$$

$$= \frac{0 + 1 + 1}{3} = \frac{2}{3} > \frac{1}{2}$$

we predict 1

An Example of Distance Weights

Use the distance weights to calculate the predicted probability and the prediction of **3NN** for D. Consider x as 1 and o as 0.



Solution



$$\text{weight} = \frac{1}{\text{Distance}}$$

$$\text{predicted prob. of 0} = \frac{\frac{1}{DC} \cdot Y_D + \frac{1}{DA} \cdot Y_A + \frac{1}{DB} \cdot Y_B}{\frac{1}{DC} + \frac{1}{DA} + \frac{1}{DB}}$$

$$= \frac{1 \cdot 0 + \frac{1}{3} \cdot 1 + \frac{1}{4} \cdot 1}{1 + \frac{1}{3} + \frac{1}{4}}$$

$$= 0.368 < \frac{1}{2}$$

predic 0

K-Nearest Neighbor: Properties

- ▶ What's nice: Simple and intuitive; easily implementable
- ▶ What's NOT nice:
 - ▶ Computationally expensive.
 - ▶ Perform not well in higher dimension data, i.e. data with many columns.

